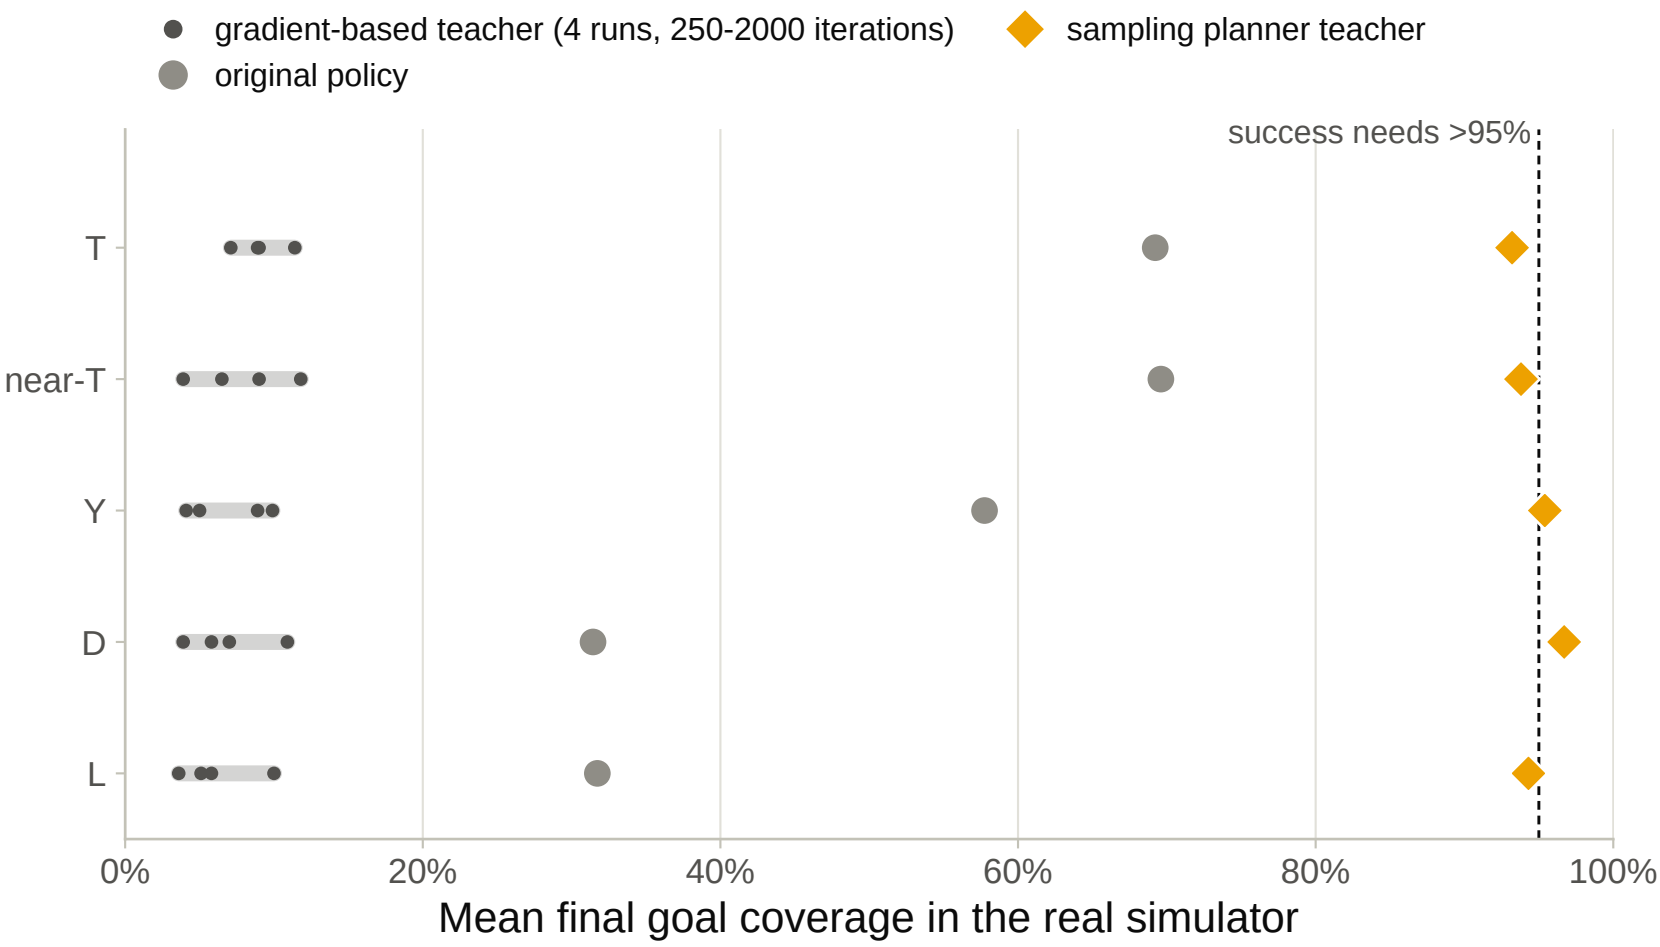

Training a teacher network by backpropagating through the learned world model (SuperMimic-style BPTT): it learned to approach the block but never pushed it anywhere (0 of 20 successes per shape in every run). The sampling planner in the same model solves most D, Y and L scenes. Data: early_attempts.csv (LEARNING_EXPERIMENT_LOG.md section 6, LEARNING_FINDINGS.md section 5); original: final_tests_by_shape.csv (D040 pool).