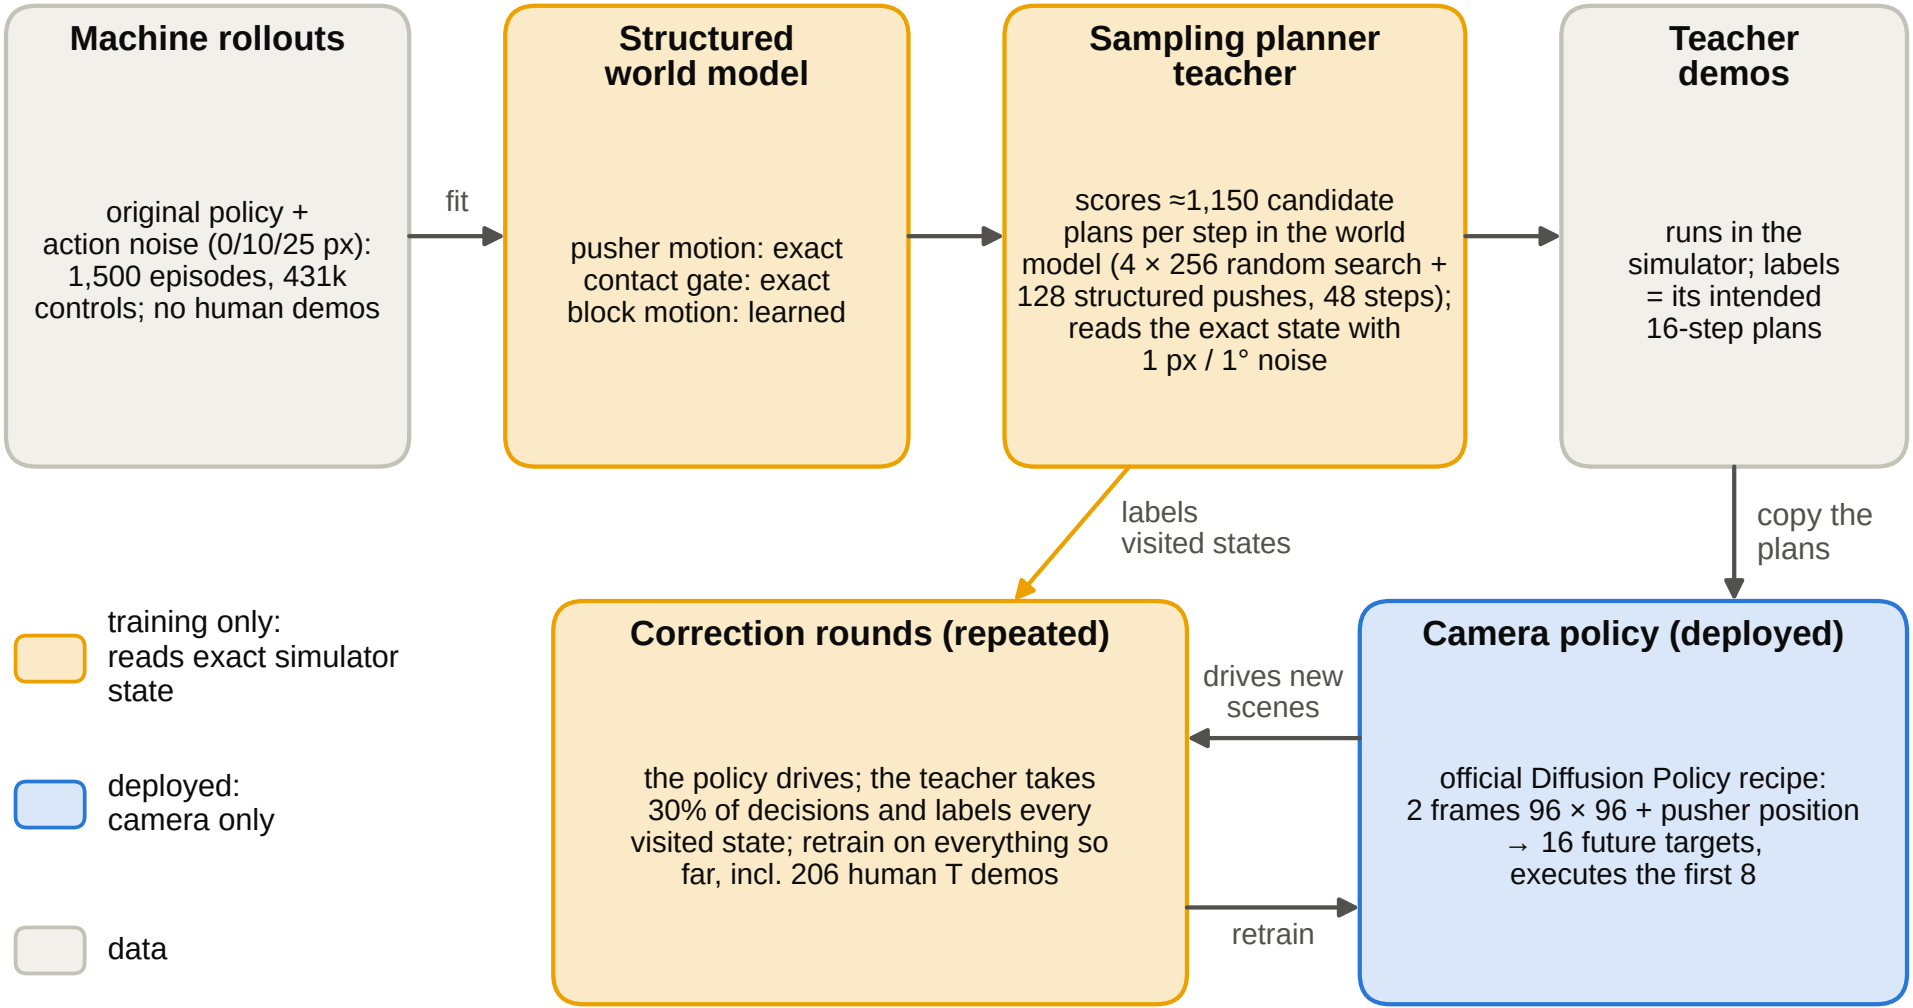
Machine rollouts
original policy + action noise (0/10/25 px): 1,500 episodes, 431k controls; no human demos
fit
Structured world model
pusher motion: exact
contact gate: exact
block motion: learned
Sampling planner teacher
scores ≈1,150 candidate plans per step in the world model (4 × 256 random search + 128 structured pushes, 48 steps); reads the exact state with 1 px / 1° noise
Teacher demos
runs in the simulator; labels = its intended 16-step plans
labels visited states
copy the plans
Correction rounds (repeated)
the policy drives; the teacher takes 30% of decisions and labels every visited state; retrain on everything so far, incl. 206 human T demos
drives new scenes
retrain
Camera policy (deployed)
official Diffusion Policy recipe: 2 frames 96 × 96 + pusher position → 16 future targets, executes the first 8
training only: reads exact simulator state
deployed: camera only
data