Left: an upper bound that uses the real simulator at decision time, not a policy (LEARNING_FINDINGS.md section 3). Right: students imitating its choices stay inside the original's run-to-run range; D and L stay at 0. The original's own plans are too similar to contain new behaviour. Data: early_attempts.csv.