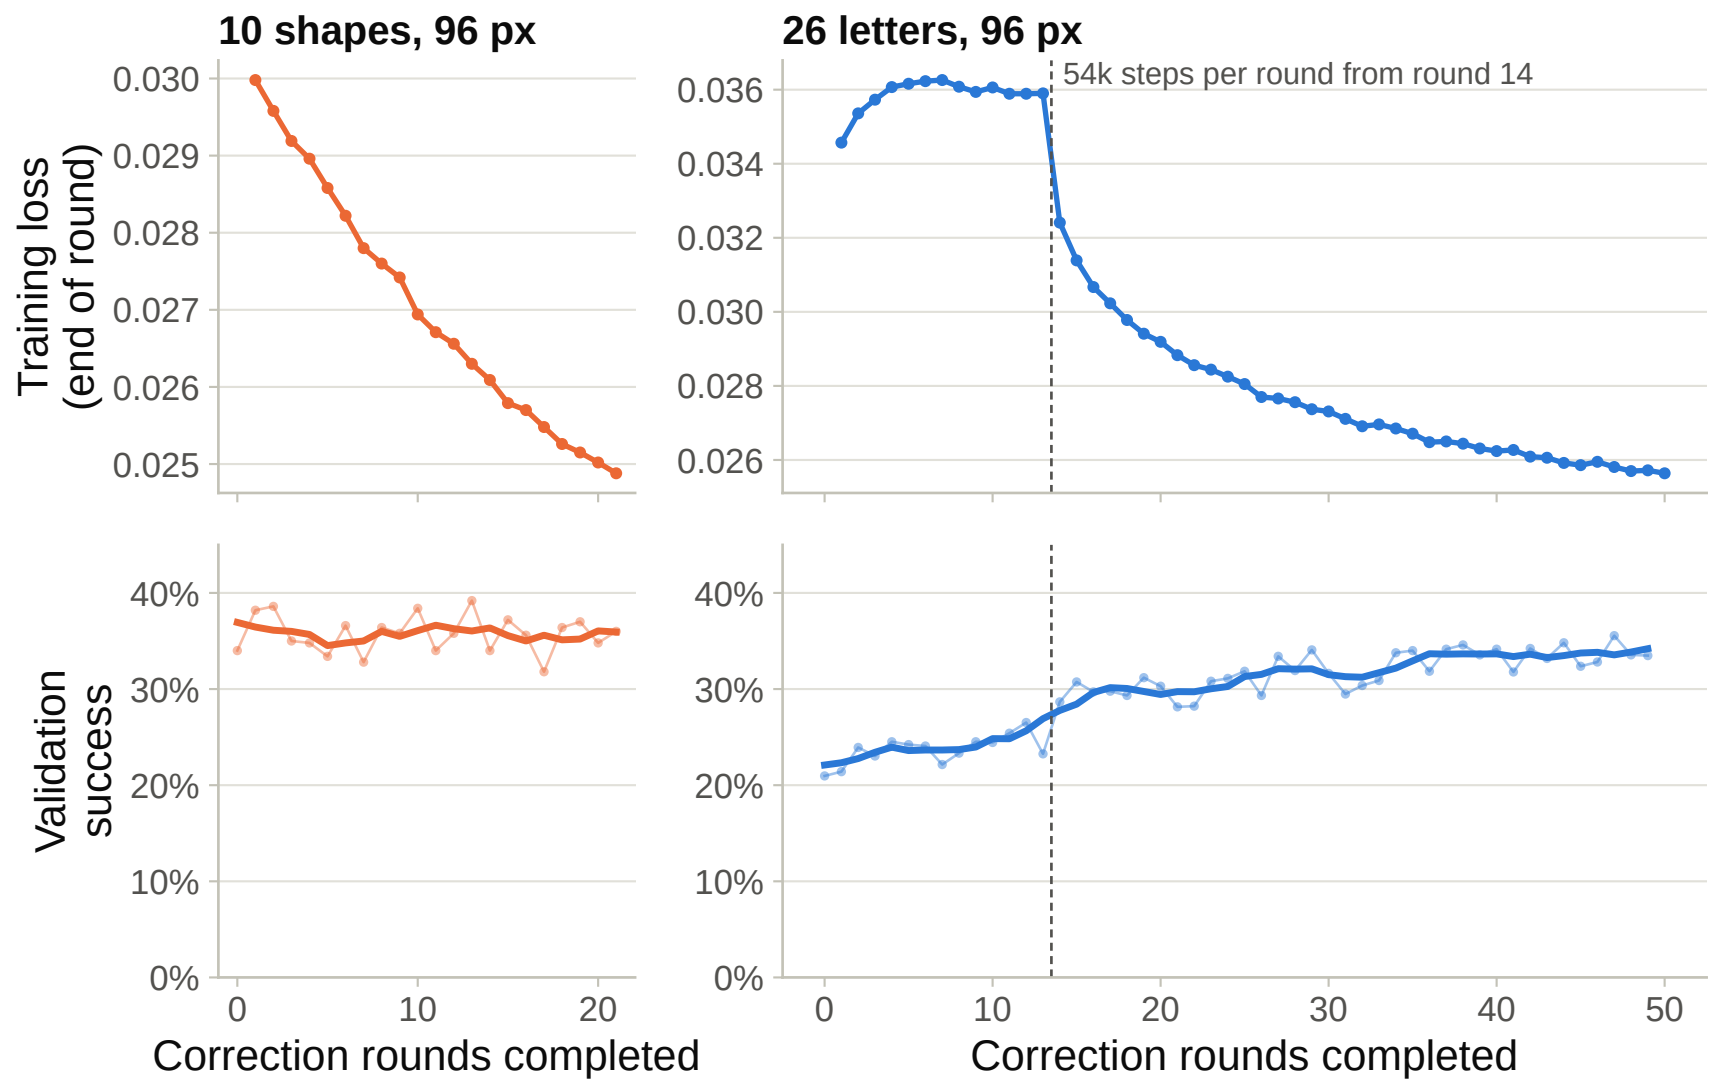

Top: mean training loss over the last 1000 steps of each round (round 0, a 30k-step retraining from the previous stage's student, and rounds that stopped early are left out). Bottom: validation success (faint) and its 5-round rolling mean. The loss keeps falling while validation stops improving: the student fits its training data better without acting better on new scenes. Data: training_rounds.csv, validation_rounds.csv.